

Private Again: AI Agents Restore Anonymity—Foreclosing Discrimination and Its Proof

Anirban Mukherjee
Hannah Hanwen Chang

Date: July 23, 2026
Forthcoming, *Stetson Law Review Forum* (Summer 2026)

Anirban Mukherjee (anirban@avyayamholdings.com) is Principal at Avyayam Holdings. Hannah H. Chang (hannahchang@smu.edu.sg; corresponding author) is Associate Professor of Marketing at the Lee Kong Chian School of Business, Singapore Management University. This research was supported by the Ministry of Education (MOE), Singapore, under its Academic Research Fund (AcRF) Tier 2 Grant, No. MOE-T2EP40124-0005.

Abstract

AI agents can transact online on behalf of a human principal—browsing, paying, receiving, and reviewing—without linking a transaction to a principal. That architecture starves algorithmic discrimination of its inputs—identity, purchase history, location history, behavioral traces, and demographic proxies—but also forecloses its proof. Disparate-treatment needs comparators; disparate-impact needs protected-class baselines; and *Iqbal*-era pleading needs specific factual allegations—doctrinal predicates that anonymous transactions never generate. The effects fall asymmetrically: those most vulnerable to discrimination are least able to afford the shield and, when harms remain, least able to prove them. The challenge for the law shifts from detecting and remedying algorithmic discrimination to governing agent-mediated anonymity as civil rights infrastructure: ensuring access to privacy-preserving agents, regulating abuse without forced identification, and deciding whether retailers may refuse to deal with agents at all.

Keywords: Privacy, Anonymity, Civil Rights, Artificial Intelligence, Antidiscrimination Law, Algorithmic Discrimination.

JEL codes: K38, K24, O33, L86.

INTRODUCTION

“Nothing was your own except the few cubic centimetres inside your skull.”

—George Orwell, *Nineteen Eighty-Four* (1949)

Each wave of technological change arrives with the same dread: that it will hollow out the private sphere.¹ The printing press prompted fears that ideas would be surveilled at the point of publication.² The telegraph and the telephone put voices into wires that, it was imagined, someone was always listening to.³ The industrial city was an alienating nightmare to those who remembered the village.⁴ The credit card replaced the unmarked handshake of cash with a permanent, centralized ledger of every purchase.⁵ Orwell’s image of the last redoubt—those few cubic centimetres inside the skull—is the most elegant expression of an anxiety that is, by now, several centuries old.⁶ In each iteration, the fear is the same: the new machine sees too much, remembers too well, and shares too widely.⁷

Artificial intelligence, however, holds the potential for an inversion: a machine that lets its user be seen less, remembered less, and traced less. As cryptocurrency has shown, the digital age need not end anonymous exchange.⁸ AI agents can extend this principle to a commercial transaction as a whole: an AI agent can serve as an ephemeral intermediary, indistinguishable from any other retail-platform user, that browses, pays, receives, and reviews for a human principal without revealing whom it represents.⁹ It can pay in cryptocurrency or with a single-use card, receive

¹See generally DANIEL J. SOLOVE, UNDERSTANDING PRIVACY 4–5 (2008) (tracing recurring privacy anxieties across successive information technologies).

²See ELIZABETH L. EISENSTEIN, THE PRINTING PRESS AS AN AGENT OF CHANGE 346–48 (1979) (describing the Archbishop of Mainz’s 1485 licensing edict and subsequent papal censorship decrees prompted by the printing press).

³See *Olmstead v. United States*, 277 U.S. 438, 473 (1928) (Brandeis, J., dissenting) (warning that government wiretapping could obtain “what is whispered in the closet”); TOM STANDAGE, THE VICTORIAN INTERNET 110–11 (1998) (describing widespread anxiety that telegraph operators and intermediaries could read every private message in transit).

⁴See RAYMOND WILLIAMS, THE COUNTRY AND THE CITY 1–12 (1973) (tracing the recurring literary contrast between an idealized rural past and a corrupt, alienating urban present).

⁵See JOSH LAUER, CREDITWORTHY: A HISTORY OF CONSUMER SURVEILLANCE AND FINANCIAL IDENTITY IN AMERICA 1–2, 184–86 (2017) (observing that each card payment enters “an invisible surveillance network” that “records the details of our transaction,” and tracing the surveillance infrastructure the cashless society required).

⁶GEORGE ORWELL, NINETEEN EIGHTY-FOUR pt. I, ch. 2 (1949) (“Nothing was your own except the few cubic centimetres inside your skull.”).

⁷See Solove, *supra* note 1, at 4–5.

⁸Bitcoin and its successors record transactions on a public ledger without necessarily linking them to real-world identities. See Satoshi Nakamoto, *Bitcoin: A Peer-to-Peer Electronic Cash System* 6 (Oct. 31, 2008) (unpublished manuscript), <https://bitcoin.org/bitcoin.pdf> [<https://perma.cc/QY6S-5EMX>] (describing a system where “[t]he public can see that someone is sending an amount to someone else, but without information linking the transaction to anyone”); PRIMAVERA DE FILIPPI & AARON WRIGHT, BLOCKCHAIN AND THE LAW: THE RULE OF CODE 20–21 (2018).

⁹See Yonadav Shavit et al., *Practices for Governing Agentic AI Systems* 4–9 (OpenAI Research Paper, Dec. 14, 2023), <https://cdn.openai.com/papers/practices-for-governing-agentic-ai-systems.pdf> [<https://perma.cc/58AG-D7GA>] (defining agentic AI systems as those that “pursue complex goals with limited direct supervision”); Noam Kolt, *Governing AI Agents*, 101 NOTRE DAME L. REV. (forthcoming 2026) (draft manuscript at 8–10), https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4772956 [<https://perma.cc/5RL9-JVFL>] (analyzing AI agents through the lens of agency law because they autonomously act on behalf of human principals).

goods at an anonymized address, and leave no readily accessible record linking the transaction to the principal at the point of sale.¹⁰ Anonymity can be restored—not by retreating from technology, but through it.

The civil rights consequences run in both directions. The architecture of algorithmic discrimination—the profiles that tie each transaction to an identifiable buyer and the models that infer protected traits from the accumulated record¹¹—collapses when the buyer is a shell: a party that cannot identify its counterparty cannot proxy-discriminate.¹² Yet that same collapse also defeats the evidentiary infrastructure on which antidiscrimination enforcement depends. Disparate-treatment needs comparators; disparate-impact needs protected-class baselines; and *Iqbal*-era pleading needs specific factual allegations.¹³ Anonymous transactions do not generate these doctrinal predicates. The mechanism that forecloses discrimination also forecloses its proof.

The challenge for the law shifts from detecting and remedying algorithmic discrimination to governing the anonymity that forecloses both. Part I of this Article traces how AI agents starve discrimination of its inputs, proof of its evidence, and pleading of its facts. Part II takes up three questions in turn—whether to treat agent-mediated anonymity as civil rights infrastructure, how to regulate abuse without forcing identification, and whether retailers may refuse to deal with agents at all.

I. CIVIL RIGHTS STARVED

Civil rights law rests on an assumption so foundational it is rarely stated: that the parties to a transaction are identifiable, whether employer and applicant, landlord and tenant, or—as here—seller and buyer.¹⁴ Discrimination, as the law understands it, is something a seller does to a buyer the seller has categorized.¹⁵ Remedy is something a court imposes after a plaintiff demonstrates that she was treated differently from a similarly situated person outside her protected class.¹⁶ Both presuppose that the seller knows the buyer’s identity and that the seller’s knowledge can be

¹⁰ See *supra* note 8.

¹¹ See *infra* Part I.A (tracing the mechanisms and the data each consumes); on the composition and durability of the profiles, see *infra* note 17; on the inference of protected traits from facially neutral data, see *infra* notes 22, 25 and accompanying text.

¹² Proxy discrimination is the use of a facially neutral trait as a stand-in for a legally protected one. The practice was historically deliberate, a means of discriminating while evading antidiscrimination law, but it need not be: a model that mines neutral data for predictive power can settle on proxies for protected traits simply because they predict. See Anya E.R. Prince & Daniel Schwarcz, *Proxy Discrimination in the Age of Artificial Intelligence and Big Data*, 105 IOWA L. REV. 1257, 1267 (2020) (“Proxy discrimination occurs when a facially-neutral trait is utilized as a stand-in—or proxy—for a prohibited trait.”).

¹³ *Ashcroft v. Iqbal*, 556 U.S. 662, 678 (2009).

¹⁴ See Samuel R. Bagenstos, *The Structural Turn and the Limits of Antidiscrimination Law*, 94 CALIF. L. REV. 1, 10–14 (2006) (analyzing antidiscrimination law’s dependence on identifying perpetrators and victims of discriminatory conduct as a precondition for liability).

¹⁵ See *Int’l Bhd. of Teamsters v. United States*, 431 U.S. 324, 335 n.15 (1977) (defining disparate treatment: the employer “simply treats some people less favorably than others because of their race, color, religion, sex, or national origin,” and “[p]roof of discriminatory motive is critical”).

¹⁶ *McDonnell Douglas Corp. v. Green*, 411 U.S. 792, 807 (1973) (explaining that a court must order a prompt and appropriate remedy after finding unlawful discrimination).

established.

This assumption is unremarkable in the traditional marketplace. Consider a reader who wishes to buy a book that carries some social or professional risk: a political memoir, a medical reference, or a work on a stigmatized topic. In the ordinary online marketplace, the purchase leaves a trail: the search query ties to an account, the order to a name and billing address, the payment to a credit card, and the shipment to a home.

Each trace is individually modest. But stitched together, they form a durable profile, one that re-identifies the reader across sites and visits. Subsequent browsing surfaces the same title or similar titles across unrelated sites.¹⁷ Prices, goods, and services follow the profile, as do advertisements. Different users with different profiles receive different treatment.

An AI agent, however, alters the picture. It can present as an ordinary user through an account the agent controls—rather than one tied to the principal’s identity—pay with a single-use virtual card or prefunded cryptocurrency, and direct shipment to a parcel locker the principal retrieves without showing identification linked to the purchase.¹⁸ The retailer sees a transaction indistinguishable from any other but cannot link that user to the principal or to the principal’s other transactions.¹⁹

The consequences are unusual precisely because they cut both ways. If a retailer cannot identify the individual behind a transaction, it cannot proxy-discriminate. It cannot price by zip code, steer by inferred demographics, or target by protected class, for there is no data from which to infer when the counterparty is an anonymous shell.²⁰ The architecture of algorithmic discrimination collapses. Anonymity, long thought a threat to civil rights, becomes their champion.

Yet the same collapse dismantles civil rights law’s remedial architecture. Antidiscrimination enforcement depends on showing that a protected class was treated differently, which presupposes identifying who was treated, by whom, and how.²¹ Affirmative remedies depend on identifying their beneficiaries. The mechanism that forecloses discrimination forecloses its proof and its

¹⁷ See JOSEPH TUROW, *THE DAILY YOU* 4–5 (2011). Several classes of identifier are used to track a user across sites: cookies and other stored identifiers, IP addresses, and device fingerprints derived from properties such as browser configuration, installed fonts, and screen resolution. The profile these identifiers assemble is durable in two senses: it persists, retained and traded by data brokers and ad networks across sessions and years; and it accretes, each transaction refining the record of identity, behavior, and inferred preference. See, e.g., Peter Eckersley, *How Unique Is Your Web Browser?*, in *PRIVACY ENHANCING TECHNOLOGIES* 1, 2–5 (Mikhail J. Atallah & Nicholas J. Hopper eds., 2010); Gunes Acar et al., *The Web Never Forgets: Persistent Tracking Mechanisms in the Wild*, in *PROCEEDINGS OF THE 2014 ACM SIGSAC CONFERENCE ON COMPUTER AND COMMUNICATIONS SECURITY* 674, 675–76, 682–83 (2014); Steven Englehardt & Arvind Narayanan, *Online Tracking: A 1-million-site Measurement and Analysis*, in *PROCEEDINGS OF THE 2016 ACM SIGSAC CONFERENCE ON COMPUTER AND COMMUNICATIONS SECURITY* 1388 (2016).

¹⁸ See *supra* note 9; *infra* note 19.

¹⁹ Each of the agent’s steps has an analogue already in wide use: privacy-preserving browsing for tracking avoidance, virtual cards for payment, parcel lockers for receipt, and temporary email for communication. See, e.g., *Tor Browser*, TOR PROJECT, <https://www.torproject.org> [<https://perma.cc/S9QJ-Y2J9>] (last visited July 4, 2026); *Privacy — Seamless & Secure Online Card Payments*, PRIVACY.COM, <https://privacy.com> [<https://perma.cc/ZML2-P94W>] (last visited July 4, 2026); *Amazon Hub Locker*, AMAZON, <https://www.amazon.com/b?node=6442600011> [<https://perma.cc/WU2H-A7ZQ>] (last visited July 4, 2026); *Proton*, *SimpleLogin*, <https://simplelogin.io> [<https://perma.cc/QH3L-NCZF>] (last visited July 4, 2026).

²⁰ See *infra* notes 23, 24, 25, and 26.

²¹ See, e.g., *McDonnell Douglas*, 411 U.S. at 800–02 (setting forth the prima facie elements required to establish a Title VII discrimination claim).

remedy.

A. Discrimination Starved of Its Inputs

The discrimination that civil rights law confronts in modern commerce is principally statistical and algorithmic²²: price discrimination keyed to zip-code models,²³ product steering driven by personalization signals,²⁴ credit scoring that relies on proxies for race and age,²⁵ and ad targeting built on behavioral aggregation.²⁶ Each mechanism consumes the same input—demographic or behavioral data tied to an identifiable buyer—producing discrimination as a byproduct. A retailer cannot steer products away from a protected class without identifying which buyers belong to that class. A lender cannot price a loan by zip code without geolocating the applicant. An ad network cannot target by inferred race if the inference has no profile to land on.

When the counterparty is an AI agent, the inputs are gone. The agent does not disclose the principal’s location, inferred demographics, purchase history, or protected class. Each session begins without history; each session ends without persistence. The agent is a shell: no stable identity to profile, no durable link to the human behind it. Proxy discrimination, a dominant modern form, collapses.²⁷ The retailer cannot engage in it. The algorithm has no signal to fit. The discrimination machinery, deliberate or inadvertent, is starved of its substrate.

²² See Solon Barocas & Andrew D. Selbst, *Big Data’s Disparate Impact*, 104 CALIF. L. REV. 671, 673–77 (2016) (describing how contemporary discrimination arises from facially neutral data-mining practices); Pauline T. Kim, *Data-Driven Discrimination at Work*, 58 WM. & MARY L. REV. 857, 860–61, 890–92 (2017) (characterizing algorithmic “classification bias” as a distinct form of modern workplace discrimination).

²³ Jennifer Valentino-DeVries, Jeremy Singer-Vine & Ashkan Soltani, *Websites Vary Prices, Deals Based on Users’ Information*, WALL ST. J., Dec. 24, 2012, at A1 (documenting Staples.com’s practice of varying prices based on a user’s inferred location and proximity to competitor stores); Ryan Calo, *Digital Market Manipulation*, 82 GEO. WASH. L. REV. 995, 1002–04 (2014) (theorizing personalized and geographically targeted pricing as a form of algorithmic consumer harm).

²⁴ See, e.g., Anikó Hannák et al., *Measuring Price Discrimination and Steering on E-commerce Web Sites*, in PROC. 2014 ACM INTERNET MEASUREMENT CONF. 305, 308–11 (2014).

²⁵ See, e.g., Robert Bartlett et al., *Consumer-Lending Discrimination in the FinTech Era*, 143 J. FIN. ECON. 30, 42–44 (2022) (finding that algorithmic lenders charge Latinx and Black borrowers higher interest rates than otherwise-similar white borrowers, even when no human loan officer is involved); Prince & Schwarcz, *supra* note 12, at 1273–76.

²⁶ See, e.g., Latanya Sweeney, *Discrimination in Online Ad Delivery*, 56 COMM. ACM 44 (2013); Muhammad Ali et al., *Discrimination Through Optimization: How Facebook’s Ad Delivery Can Lead to Biased Outcomes*, 3 PROC. ACM HUM.-COMPUT. INTERACTION, art. 199, at 12–15 (Nov. 2019); Charge of Discrimination, Sec’y, U.S. Dep’t of Hous. & Urb. Dev. v. Facebook, Inc., FHEO No. 01-18-0323-8 (Mar. 28, 2019), https://archives.hud.gov/news/2019/HUD_v_Facebook.pdf [<https://perma.cc/3MJC-MGJJ>].

²⁷ See Prince & Schwarcz, *supra* note 12, at 1275–78. Prince and Schwarcz show that depriving an algorithm of a protected characteristic does not prevent proxy discrimination, because a sophisticated model “will identify and use any available data that even partially proxies for th[at] information”—geolocation, spending patterns, family medical history. *Id.* at 1278; see *id.* at 1276. That mechanism, however, presupposes the presence of proxies. Under agent-mediated anonymity, the agent’s shell withholds not merely the protected characteristic but all possible proxies. With nothing available from which to reconstruct, the model can do no better than to treat every anonymous counterparty alike—it cannot discriminate. This protection is structural rather than behavioral: agent-mediated anonymity removes the capacity to proxy-discriminate rather than relying on the seller’s forbearance. The point holds, though, only insofar as the agent is a genuine shell. Device fingerprints, payment metadata, a shipping address, or cross-session linkage might defeat the anonymity the shell provides, reintroducing the linked signals on which reconstruction depends.

B. Proof Starved of Its Evidence

The same mechanism that defeats discrimination defeats the enforcement apparatus built to detect and punish it. Disparate-treatment doctrine requires the plaintiff to show she was treated worse than a similarly situated comparator outside her protected class.²⁸ That showing requires identifying herself (the plaintiff), the defendant's knowledge (what the defendant could have known about her), and a comparator (someone similarly situated but outside the class, who was treated differently). Each step depends on records that AI agents do not produce. The plaintiff identifies herself only by stepping out of the anonymity regime and disclosing the information the agent was built to conceal; even then, the defendant may have no record of who she was during the transaction. The defendant lacks the requisite knowledge because the transactional interface provides nothing to know. And the comparator is an agent, indistinguishable from the plaintiff's own agent by design.

Disparate-impact doctrine fares no better. *Griggs v. Duke Power Co.* and its progeny require proof that a facially neutral practice produces statistically significant differential effects on a protected class.²⁹ That proof requires identifying the protected class within the defendant's customer base and measuring outcomes against that identification. If every customer is an agent, there is no protected class to identify, no demographic baseline to compare against, and no statistical machinery to deploy. The doctrine does not fail because the practice is fair. It fails because the data to test fairness does not exist. Discovery does not save the case: records the defendant does not keep cannot be produced, and a retailer who sees only agents has no customer-level demographics to disgorge. Expert testimony encounters the same wall. An expert retained to show that a retailer's checkout flow produces differential outcomes across protected classes needs the ground truth of user identity; a retailer whose records contain only agents has no such ground truth to provide. The expert is left testifying about a pattern in data that, under an anonymity regime, does not exist.

C. Pleading Starved of Its Facts

The problem with enforcement is not only that the evidentiary record is thin at trial. It is that the record does not exist at the threshold at which enforcement begins. *Iqbal* and *Twombly* elevated the pleading standard from "conceivable" to "plausible," requiring a complaint to allege concrete facts from which unlawful conduct can be reasonably inferred.³⁰ A plaintiff must come to the courthouse with facts sufficient to make discrimination a *plausible*, not merely conceivable, explanation for the defendant's conduct.³¹

Those facts come from three sources. The plaintiff may have *direct evidence*: a memo, an email, a policy, or a remark revealing bias.³² She may have *comparative evidence*: knowledge

²⁸ See *McDonnell Douglas Corp. v. Green*, 411 U.S. 792, 802–04 (1973) (establishing the burden-shifting framework for disparate-treatment claims).

²⁹ *Griggs v. Duke Power Co.*, 401 U.S. 424, 431 (1971); see also *Tex. Dep't of Hous. & Cmty. Affairs v. Inclusive Cmty. Project, Inc.*, 576 U.S. 519, 540 (2015) (affirming disparate-impact liability under the Fair Housing Act).

³⁰ *Ashcroft v. Iqbal*, 556 U.S. 662, 678 (2009); *Bell Atl. Corp. v. Twombly*, 550 U.S. 544, 570 (2007).

³¹ See cases cited *supra* note 30.

³² See *Price Waterhouse v. Hopkins*, 490 U.S. 228, 250–52 (1989) (plurality opinion) (recognizing that a plaintiff may

that similarly situated comparators outside her protected class were treated differently.³³ Or she may have *statistical evidence*: patterns in a defendant’s behavior revealed by prior litigation, investigative journalism, regulatory disclosure, or organized scrutiny.³⁴ In a world of identifiable customers, at least one of these is usually available. The architecture of discrimination enforcement depends on that availability.

Agent-mediated anonymity removes all three sources at once. There is no direct evidence because the agent strips the demographic signal that would trigger discriminatory behavior: no email or memo can reveal bias against a class the defendant never knew it was transacting with. There is no comparative evidence because comparators are themselves anonymous agents, indistinguishable from the plaintiff’s own agent by construction. And there is no statistical evidence because no defendant, regulator, or third party can assemble demographic data about customers who present only as shells. The pleading standard, faithfully applied, results in dismissal, not because the plaintiff’s claim is implausible, but because the facts that would make *any* such claim plausible cannot exist.

The result is a doctrinal closed loop. A plaintiff cannot plead what she cannot know; she cannot know what she cannot discover; she cannot discover without first pleading. The gate locks from both sides. *Iqbal* is not malfunctioning; the Court’s concern about abusive pleading is met exactly as intended.³⁵ But *Iqbal*-era pleading rigor combined with agent-mediated anonymity forecloses the courthouse door to discrimination claims previously colorable under the permissive pleading regime.³⁶ The enforcement apparatus is not merely weakened. Its entry point is closed.³⁷

The same closure affects the affirmative side of antidiscrimination law, which seeks not only to

prove discrimination by showing that an illegitimate criterion “played a motivating part” in the employment decision, including through direct expressions of bias); Michael J. Zimmer, *The New Discrimination Law: Price Waterhouse Is Dead, Whither McDonnell Douglas?*, 53 EMORY L.J. 1887, 1888–907 (2005) (tracing the doctrinal role of direct evidence in disparate-treatment cases and the development of the “stray remarks” doctrine in the lower courts).

³³ See *McDonnell Douglas*, 411 U.S. at 804 (suggesting that a plaintiff may show pretext through evidence that “white employees involved in acts against [the employer] of comparable seriousness . . . were nevertheless retained or rehired”); Suzanne B. Goldberg, *Discrimination by Comparison*, 120 YALE L.J. 728, 732–35 (2011) (analyzing the centrality and limits of the “similarly situated” comparator requirement in antidiscrimination doctrine).

³⁴ See *Hazelwood Sch. Dist. v. United States*, 433 U.S. 299, 307–08 (1977) (reaffirming that “gross statistical disparities” can themselves “constitute prima facie proof of a pattern or practice of discrimination”); *Int’l Bhd. of Teamsters v. United States*, 431 U.S. 324, 339–40 (1977) (recognizing that “statistics . . . come in infinite variety” and may be probative of pattern-or-practice discrimination).

³⁵ See *Iqbal*, 556 U.S. at 678–79.

³⁶ See *Conley v. Gibson*, 355 U.S. 41, 45–46 (1957) (establishing the “no set of facts” pleading standard subsequently abrogated by *Bell Atl. Corp. v. Twombly*, 550 U.S. 544, 570 (2007)).

³⁷ The argument that the pleading stage, rather than the ultimate burden of proof, is the operative failure point for discrimination claims in an anonymity regime appears novel to AI. Prior literature on discrimination pleading after *Iqbal* has focused on the evidentiary difficulties faced by particular classes of plaintiffs in particular doctrinal settings, see, e.g., Joseph A. Seiner, *The Trouble with Twombly: A Proposed Pleading Standard for Employment Discrimination Cases*, 2009 U. ILL. L. REV. 1011, 1014–15; Charles A. Sullivan, *Plausibility Pleading Employment Discrimination*, 52 WM. & MARY L. REV. 1613, 1621 (2011), without considering the combinatorial problem posed when the underlying identification architecture is absent altogether. The closed-loop structure described here parallels the “intermediate copying double bind” identified in the AI copyright context: claims cannot be initiated without the very evidence that the structure of the technology prevents from existing. See Anirban Mukherjee & Hannah Hanwen Chang, *Pro Forma Copyright: AI Compaction and the Idea-Expression Impasse* 22 (May 6, 2026) (unpublished manuscript), <https://ssrn.com/abstract=6232359> [<https://perma.cc/4D9H-NTUP>].

prevent unequal treatment in the moment but also to remedy the effects of historical marginalization through targeted intervention, including affirmative action in contracting, minority business development programs, and targeted lending initiatives.³⁸ Each presupposes that the beneficiary can be identified as a member of the target group. At the point of transaction, anonymity strips that identification: an agent does not announce its principal’s race to a minority business development program; a cryptocurrency payment does not identify its source community; a minority business set-aside cannot operate when bidders are agents. At the transaction, the mechanism that prevents harm also prevents help.

II. THE ROLE OF LAW

The return of anonymity to commerce poses three questions for the law: whether anonymity should be treated as civil rights infrastructure, available to those who cannot afford it; how to permit legitimate use while foreclosing abuse; and whether retailers may refuse to transact with agents at all—and if so, what that refusal means for the anonymity this Article has described.

A. Anonymity as Civil Rights Infrastructure

Sophisticated AI agents are not free.³⁹ They require technical literacy, computational resources, and a willingness to navigate a commercial world that has not yet adapted to them.⁴⁰ Early adopters will skew toward the populations that already enjoy the best privacy tools: the wealthy, the technically fluent, and those with the time to configure their transactional lives carefully.

This is the digital divide at a new layer. Agent-mediated commerce will leave non-agent users exposed to the price discrimination, behavioral targeting, and algorithmic harms that agent-users escape. The discrimination machinery will not disappear; it will shift from targeting members of a protected class to targeting the residual population that did not, or could not, deploy the shield. The defeat of discrimination against agent-users is not the defeat of discrimination; it is its concentration on those who remain visible.

The obvious response is to sever the link between anonymity and wealth. If privacy-enhancing agents become the most effective protection against the discrimination machinery Part I described, access to those agents should not be contingent on the ability to pay. Civil rights infrastructure has addressed analogous problems before. Legal Aid exists because legal representation is expensive; the EEOC exists because individual enforcement is resource-intensive; housing counselors exist because the information asymmetries in mortgage and rental markets are severe.⁴¹ Each rests on

³⁸ See, e.g., 15 U.S.C. § 644(g)(1)(A)(iv); 15 U.S.C. § 637(a); 12 U.S.C. § 2901.

³⁹ See Jiin Kim et al., *The Cost of Dynamic Reasoning: Demystifying AI Agents and Test-Time Scaling from an AI Infrastructure Perspective*, arXiv, Jan. 7, 2026, at 1–2, arXiv:2506.04301 [cs.LG] [<https://perma.cc/3B7E-8SSE>] (finding that agentic, multi-step AI workflows impose substantial resource, latency, energy, and infrastructure costs).

⁴⁰ See Lijia Ma et al., *Learning to Adopt Generative AI*, arXiv, Feb. 15, 2026, at 2–4, arXiv:2410.19806 [cs.CY] [<https://perma.cc/6QDA-W9MP>] (arguing that AI adoption and effective use vary according to education, technological background, and the ability to learn from experience).

⁴¹ See ALAN HOUSEMAN & LINDA E. PERLE, SECURING EQUAL JUSTICE FOR ALL: A BRIEF HISTORY OF CIVIL LEGAL ASSISTANCE IN THE UNITED STATES 60 (Ctr. for L. & Soc. Pol’y, rev. ed., 2018); 42 U.S.C. § 2000e-4 (establishing the Equal Employment Opportunity Commission); 24 C.F.R. pt. 214 (2026) (HUD housing counseling program).

the same principle: where a protection is necessary to prevent a recognized form of discrimination, access to the protection should not be limited to those who can afford it.

A parallel response would treat the AI agent itself as civil rights infrastructure—publicly provided, subsidized, or required to be offered free by platform providers above a certain scale. The analogues are imperfect but not absent. Public libraries provide free internet access because connectivity is now necessary for participation in civic life. Universal service rules subsidize telecommunications in underserved areas because the telephone network became too important to be rationed by the ability to pay.⁴²

Agent-mediated anonymity could be framed in the same register if the law recognizes that what is being protected is not merely consumer preference, but the structural precondition for civil rights enforcement in a post-identification marketplace. Public libraries and universal service are analogues for access to a positive good; the closest precedent is *Gideon v. Wainwright*'s right to counsel—a protection against state power that the state itself provides.⁴³

B. Permitting Use, Foreclosing Abuse

Anonymity has always been double-edged. Bitcoin enabled private exchange, and it enabled the Silk Road marketplace and its successors.⁴⁴ End-to-end encryption protects the dissident and the predator both.⁴⁵ The history of privacy-enhancing technology is, in significant part, a history of regulators negotiating how much abuse to tolerate in exchange for how much legitimate protection. Agent-mediated commerce will replay that negotiation at a higher level of abstraction.

The distinctive difficulty is that agent-mediated anonymity is not a shield around one dimension of the transaction—value, content, identity—but around the transaction as a whole. A regulator responding to Silk Road could focus on the payment rail: Bitcoin in, fiat out, and know-your-customer (KYC) requirements at the on-ramps and off-ramps. A regulator responding to encrypted messaging could focus on the endpoint: the device, the app provider, and the metadata. A regulator responding to agent-mediated commerce faces a technology whose function is to make every point in the transaction opaque simultaneously. There is no clean rail at which identification can be mandated. The shield has no seams.

The KYC and anti-money-laundering regimes illustrate the challenge. Existing financial regulation requires identification of counterparties to certain transactions as a condition of participation in the banking system.⁴⁶ AI agents transacting in cryptocurrency or single-use virtual card num-

⁴² See 47 U.S.C. § 254 (universal service provisions added by the Telecommunications Act of 1996); 47 C.F.R. pt. 54, subpt. E (2026) (Lifeline program subsidizing telecommunications service for low-income consumers).

⁴³ See *Gideon v. Wainwright*, 372 U.S. 335, 344 (1963) (“[R]eason and reflection require us to recognize that in our adversary system of criminal justice, any person haled into court, who is too poor to hire a lawyer, cannot be assured a fair trial unless counsel is provided for him”).

⁴⁴ *United States v. Ulbricht*, 858 F.3d 71, 82–83 (2d Cir. 2017); see also Nicolas Christin, *Traveling the Silk Road: A Measurement Analysis of a Large Anonymous Online Marketplace*, in *PROC. OF THE 22D INT’L CONF. ON WORLD WIDE WEB* 213, 220–22 (2013) (estimating Silk Road’s annual revenue and transaction volume).

⁴⁵ See STEVEN LEVY, *CRYPTO: HOW THE CODE REBELS BEAT THE GOVERNMENT—SAVING PRIVACY IN THE DIGITAL AGE* 197–98 (2001).

⁴⁶ Bank Secrecy Act, 31 U.S.C. § 5318(l) (requiring financial institutions to implement customer identification programs); 31 C.F.R. § 1020.220 (2026) (bank Customer Identification Program rule requiring identity-verification procedures sufficient to form a reasonable belief that the bank knows the customer’s true identity); FIN. ACTION TASK FORCE,

bers sit uneasily with these requirements. A response that requires the agent itself to be identified defeats the architecture; one that exempts agent-mediated transactions entirely turns the agent into a laundering tool.

By making every point in the transaction opaque simultaneously, agent-mediated commerce compels the regulatory architecture to sit above or below the transaction layer rather than within it.⁴⁷ Two directions are worth considering. Above the transaction, at the gateway, agent providers might bear obligations to maintain audit trails releasable only upon judicial authorization, thereby preserving a pathway to identification when circumstances warrant. Below the transaction, at the remedial layer, victims of agent-enabled fraud might be made whole through ex ante insurance or distributive compensation rather than through retrospective investigation, which anonymity renders impractical.

C. The Retailer Ban

The private side of the market is likely to be less accommodating. Retailers have strong incentives to identify their customers. Identification enables personalized pricing, targeted upselling, profile building, resale of behavioral data, and—not least—the commercial surveillance and discrimination that Part I described as socially costly but that are, from the retailer’s perspective, profit-generating. A retailer facing a customer who presents as an agent sees a customer whose transactional value has been narrowed to the margin on the immediate sale. Everything else the retailer would normally extract is gone.

The natural response is refusal. A retailer can amend its Terms of Service to ban agent-mediated transactions, require proof of human identity before completing a purchase, implement CAPTCHA-like challenges designed to defeat automation, or simply decline to serve any account it suspects of being an agent. Each restores the retailer’s profile-building capacity by forcing the customer to present as a human with a persistent identity. Under current law, each is presumptively permissible: private actors generally set the terms of their own commerce, and the idea that a retailer might be *obligated* to serve anonymous customers cuts sharply against the background presumption of freedom of contract.⁴⁸

But freedom of contract is not absolute. Public-accommodations law already limits a retailer’s

International Standards on Combating Money Laundering and the Financing of Terrorism & Proliferation: The FATF Recommendations, Recommendation 10 (June 2026), <https://www.fatf-gafi.org/en/publications/Fatfrecommendations/Fatf-recommendations.html> [<https://perma.cc/YK6S-D5YJ>]; FIN. CRIMES ENF’T NETWORK, FIN-2019-G001, *Application of FinCEN’s Regulations to Certain Business Models Involving Convertible Virtual Currencies*, at 15–17 (May 9, 2019), <https://www.fincen.gov/sites/default/files/2019-05/FinCEN%20Guidance%20CVC%20FINAL%20508.pdf> [<https://perma.cc/ZHM6-RTQ5>].

⁴⁷Prior scholarship on cryptocurrency regulation, KYC/AML obligations, and platform governance has generally assumed that the transaction remains the operative regulatory unit. See, e.g., Sarah Jane Hughes & Stephen T. Middlebrook, *Advancing a Framework for Regulating Cryptocurrency Payments Intermediaries*, 32 YALE J. ON REGUL. 495, 549–50 (2015); Kate Klonick, *The New Governors: The People, Rules, and Processes Governing Online Speech*, 131 HARV. L. REV. 1598, 1635 (2018). Agent-mediated commerce, however, forecloses that assumption.

⁴⁸See generally P.S. ATIYAH, *THE RISE AND FALL OF FREEDOM OF CONTRACT* (1979) (tracing the rise and decline of contractual autonomy as an organizing principle of Anglo-American private law); LAWRENCE M. FRIEDMAN, *CONTRACT LAW IN AMERICA: A SOCIAL AND ECONOMIC CASE STUDY* (1965) (examining the social and economic forces shaping American contract doctrine).

freedom to refuse service on protected-class grounds.⁴⁹ If agent-mediated anonymity is, as Part I argued, an effective protection against algorithmic discrimination, then a retailer's refusal to serve agents is a refusal to serve customers through a channel that prevents discrimination. The functional consequence is the restoration of the retailer's capacity to discriminate. The law's background presumption—that private actors set their own terms—runs up against the law's foreground commitment—that private actors may not set those terms in ways that enable discrimination.

Framed narrowly, that is a business decision—a private actor setting the terms of its own commerce. Framed more honestly, it is a demand for identification as a condition of participation in the commercial sphere. That demand has its own history in American civil rights law, and history is not on its side. Pass laws in colonial and antebellum America required free Black people to carry papers establishing their right to be where they were; a person who could not produce them was presumed enslaved and could be jailed, hired out, or sold.⁵⁰ Voter identification debates continue to turn on the disparate impact of identification requirements on the same populations the civil rights regime is designed to protect; where courts have found such requirements discriminatory, by result or by purpose, they have held them unlawful.⁵¹ A commercial requirement that every purchaser be identifiable is structurally analogous to these earlier regimes of forced identification: imposed each time in the name of security, operating each time as an instrument of exclusion.⁵²

⁴⁹42 U.S.C. § 2000a; see *Heart of Atlanta Motel, Inc. v. United States*, 379 U.S. 241, 258–61 (1964); cf. *Masterpiece Cakeshop, Ltd. v. Colo. Civil Rts. Comm'n*, 584 U.S. 617, 623–24 (2018) (addressing the tension between public-accommodations law and religious and free-speech objections to serving certain customers).

⁵⁰See IRA BERLIN, *SLAVES WITHOUT MASTERS: THE FREE NEGRO IN THE ANTEBELLUM SOUTH* 92–96, 327–40 (1974) (describing registration requirements, certificates of freedom, and pass laws imposed on free Black people throughout the antebellum South, and the jailing, hiring out, and sale of those found without papers); THOMAS D. MORRIS, *SOUTHERN SLAVERY AND THE LAW, 1619–1860*, at 22–36 (1996) (tracing the legal definitions of race and the presumption that a Black person was enslaved, leaving freedom to be affirmatively proven).

⁵¹See *Veasey v. Abbott*, 830 F.3d 216, 249–52 (5th Cir. 2016) (en banc) (holding that Texas's voter identification law had a discriminatory effect on minority voters in violation of Section 2 of the Voting Rights Act); N.C. State Conf. of the NAACP v. *McCrory*, 831 F.3d 204, 214–15 (4th Cir. 2016) (holding that North Carolina's photo-identification requirement and related voting restrictions were enacted with racially discriminatory intent in violation of Section 2 of the Voting Rights Act and the Fourteenth Amendment, and observing that the challenged provisions “target African Americans with almost surgical precision”); *Crawford v. Marion Cnty. Election Bd.*, 553 U.S. 181, 198–203 (2008) (plurality opinion) (upholding Indiana's voter-ID law but acknowledging the burden on voters who lack qualifying identification); Zoltan L. Hajnal, Nazita Lajevardi & Lindsay Nielson, *Voter Identification Laws and the Suppression of Minority Votes*, 79 J. POL. 363, 371–73 (2017). But see *Frank v. Walker*, 768 F.3d 744, 753 (7th Cir. 2014) (acknowledging findings that “document a disparate outcome” in identification possession but holding that Section 2 “does not condemn a voting practice just because it has a disparate effect on minorities”).

⁵²The structural analogy between a retailer's Terms-of-Service ban on agent-mediated transactions and historical forced-identification regimes (pass laws, voter ID requirements) appears novel to AI commerce. Prior scholarship on algorithmic discrimination has generally focused on how systems use observed or inferred traits once individuals are legible to the system, see, e.g., Kim, *supra* note 22, at 860–61; Anupam Chander, *The Racist Algorithm?*, 115 MICH. L. REV. 1023, 1028–31 (2017), rather than on whether forced identification is itself a contested civil rights question.

III. CONCLUSION

Orwell's Winston Smith thought the inside of his skull was the last uncolonized ground. This Article has argued that the perimeter of the private can, in at least one dimension of modern life, be pushed back outward again. AI agents can return to commerce a degree of anonymity unavailable in practice since cash. The civil rights implications are unusual: the mechanism of algorithmic discrimination collapses when the counterparty is a shell, as does civil rights law's detection-and-remedy architecture. Anonymity protects on one side and forecloses on the other, falling hardest on the populations the civil rights project was built to protect. Whether the law recognizes the possibility, and whether it protects that possibility against private actors incentivized to deny it, is the question next-generation commercial regulation must face.